

AVISH JADWANI

Senior Data Engineer

Austin, Texas | +1-857-544-6635 | jadwaniavish@gmail.com | [LinkedIn](#) | [GitHub](#) | [Website](#)

SUMMARY

Senior Data Engineer with 6+ years building production-grade data pipelines, warehouses, and analytics products for regulated, privacy-sensitive organizations, including pipelines sourced from HR/ATS systems such as Workday, Greenhouse, and LinkedIn. Deep hands-on expertise in SQL across distributed (Trino/Presto/Hive) and relational (Postgres) engines, Airflow orchestration, and Python integrations with APIs, S3, and SFTP sources. Proven track record standing up secure data foundations for LLM/RAG and AI-driven tools — including access controls, sensitivity classification, and evaluation/feedback datasets — and building guardrail APIs that redact, hash, and role-gate access to sensitive data models — and delivering Streamlit dashboards that make complex data accessible to non-technical business stakeholders. Experienced partnering cross-functionally with HR, Legal, Risk, and analytics teams to translate people- and business-related problems into governed, scalable data solutions.

CORE COMPETENCIES

SQL & Query Engines: Trino, Presto, Hive, Postgres (OLAP & OLTP), query optimization at petabyte scale

Data Engineering: Airflow, DBT, Spark, PySpark, Kafka, Flink, HDFS/Hadoop, Iceberg, Delta Lake

Databases & Storage: Redshift, Snowflake, BigQuery, Postgres, DynamoDB, MongoDB, Cassandra

Cloud & Infra: AWS (EC2, Lambda, Glue, S3), Ubuntu/Linux administration, SSH, Kubernetes, Docker, Terraform, Jenkins CI/CD

AI / LLM Data Layers: RAG architecture, LangChain, LlamaIndex, Vector DBs (Qdrant, Pinecone), AWS Bedrock, feedback datasets

HR / People Systems: Workday, Greenhouse, LinkedIn — pipeline design & integration into governed warehouses

Data Products: Streamlit dashboards & reporting apps, BI-ready data models, stakeholder-facing analytics

Governance & Security: Sensitivity classification, field-level redaction/hashing, audit logging, RBAC/MFA, Azure AD/OAuth2

Practices: Git version control, code review & mentoring, data modeling (keys, indexes, partitioning), Agile/SDLC

PROFESSIONAL EXPERIENCE

Snap Finance

May 2022 – Present

Senior Data Engineer

Austin, Texas

- Built the underlying data infrastructure — pipelines, access controls, and governed data models — to move an internal LLM-powered tool from prototype to production: architected a secure, **cloud-native RAG service on AWS Bedrock, FastMCP, and Qdrant**, deployed on **Kubernetes**, reducing **data inquiry tickets by 70%**.
- Led design and delivery of a centralized data warehouse (**DBT + Airflow**) that built and maintained pipelines from 10+ disparate source systems — including **Workday, Greenhouse, LinkedIn, Google, Facebook, and Salesforce** — into a unified, governed data model for BI and reporting, **cutting analysis turnaround from days to hours**.
- Optimized **SQL** across distributed (**Trino**) & relational engines to support high-frequency BI and analytics use cases, & architected a **medallion warehouse** model with **PySpark** ETL to transform nested event data into denormalized tables, **improving query performance by 40%**.
- Engineered a fault-tolerant, scalable **Spark/Iceberg** merchant attribution pipeline that reallocated over \$1M in budget, demonstrating platform reliability and design at scale.
- Designed and built a **Guardrails API** for querying critical data models containing sensitive data, **implementing sensitivity classification, audit logging, and role-based access control** to redact, hash, or deny access to protected fields based on requester permissions.
- Built and maintained secure **SFTP ingestion pipelines** to pull financial data from external sources such as NetSuite, integrating them into the centralized warehouse alongside API- and S3-based sources.
- Set organization-wide **DBT modeling standards**, led adoption of automated linting and pre-commit hooks via Git, and **mentored engineers** across distributed teams to enforce secure, consistent, high-quality code.

State Auto Insurance

Jul 2021 – May 2022

Data Engineer

Columbus, Ohio

- Implemented secure data integrations using **OAuth2** and **Azure Active Directory** with role-based access control (**RBAC**) and multi-factor authentication (**MFA**) to protect sensitive, regulated data — directly applicable to permissioning and access-control requirements for sensitive employee data.
- Triaged, debugged, and **optimized SQL queries** against petabyte-scale distributed claims and policy datasets, cutting actuarial report execution time from hours to minutes while preserving data quality and system reliability.
- Built and led adoption of an automated **CI/CD pipeline (Jenkins, Docker, Kubernetes)** for data platform deployments in a regulated, privacy-sensitive environment, improving deployment efficiency 3x and **reducing production errors 45%**.

Marsh McLennan

Jan 2020 – Sep 2020

Data Science Intern

San Francisco, CA

- Partnered with forecasting and analytics teams to translate business requirements into a scalable risk-assessment data product, combining a TensorFlow model with a **Streamlit** application — a lightweight frontend framework for surfacing insights to non-technical stakeholders.
- Built and scaled a custom Python feature-engineering library integrated with **Apache Spark** to process terabytes of policy and claims data, accelerating model development by 25%.
- Documented **SLAs** and technical specifications for datasets, establishing clear data-quality guidelines across areas of ownership.

Cognizant Technology Solutions

Jul 2016 – Jan 2018

Software Engineer

Bangalore, India

- Developed secure, production-quality code for user-facing features in **React.js, Redux, and JavaScript**, and built reusable components in Java to streamline development across teams.
- Conducted code reviews focused on coding standards, best practices, and performance for peer engineers.

EDUCATION

Northeastern University

2021

Master of Science, Data Analytics Engineering — GPA: 3.95/4

VIT University

2016

Bachelor of Technology, Mechanical Engineering — GPA: 3.75/4